

< Prev (<https://www.ops-class.org/slides/2017-02-08-forkandfilehandles/>) (deck.html) Next > (<https://www.ops-class.org/slides/2017-02-13-synchprimitives/>)

fork() and pipe()

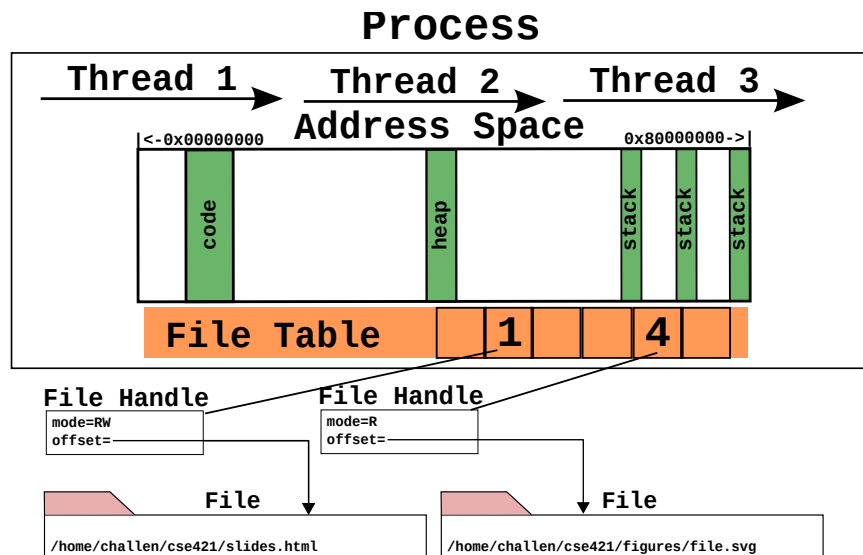
Process Creation

Where do processes come from?

fork() # create a new process

fork() is the UNIX system call that creates a new process.

- fork() creates a new process that is a **copy** of the calling process.
- After fork() we refer to the caller as the **parent** and the newly-created process as the **child**. This relationship enables certain capabilities.



fork() Semantics

- Generally fork() tries to make an **exact** copy of the calling process.
 - Recent version of UNIX have relaxed this requirement and there are now many flavors of fork() that copy different amounts of state and are suitable for different purposes.
 - For the purposes of this class, ignore them.
- Threads are a notable exception!

fork() Against Threads

- Single-threaded fork() has reliable semantics because the **only thread** the processes had is the one that called fork().
 - So nothing else is happening while we complete the system call.

- Multi-threaded fork() creates a host of problems that many systems choose to ignore.
 - Linux will only copy state for the thread that called fork().

Multi-Threaded fork()

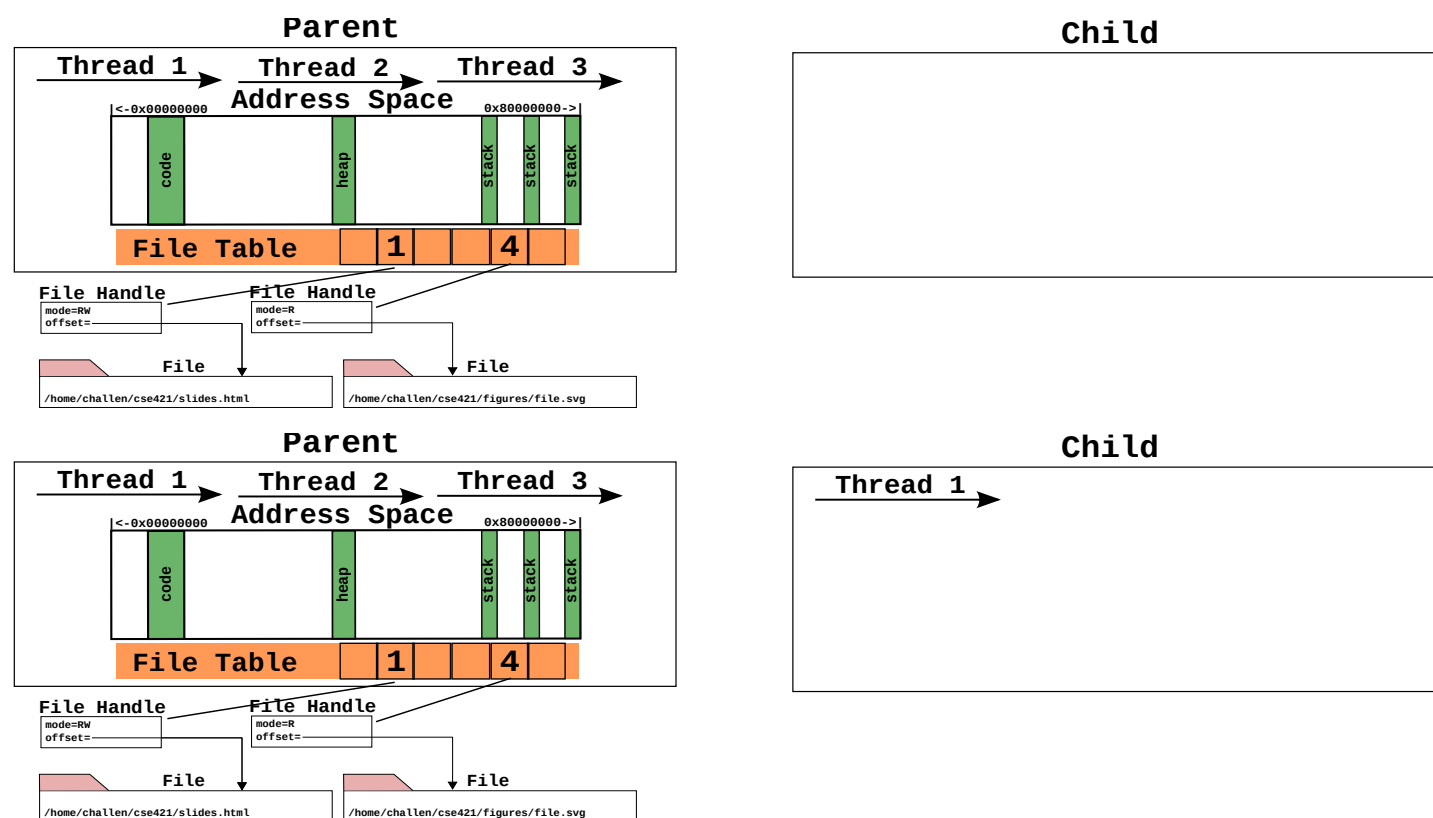
There are two major problems with multi-threaded fork()




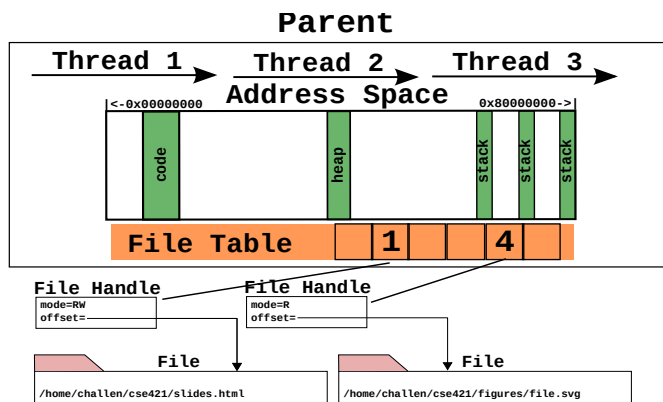

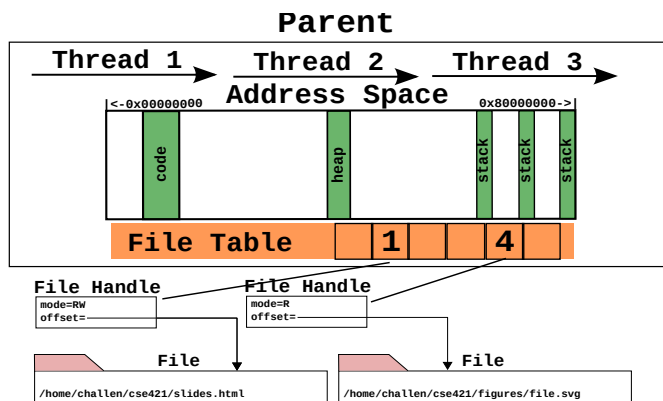
1. Another thread could be blocked in the middle of doing something (uniprocessor systems), or
2. another thread could be **actually** doing something (multiprocessor systems).

This ends up being a big mess. Let's just copy the calling thread.

fork()

1. fork() copies one thread—the caller.
2. fork() copies the address space.
3. fork() copies the process file table.





After fork()

```
returnCode = fork();
if (returnCode == 0) {
    # I am the child.
} else {
    # I am the parent.
}
```

- The child thread returns executing at the exact same point that its parent called fork().
 - With one exception: fork() returns **twice**, the PID to the parent and 0 to the child.
- All contents of memory in the parent and child are identical.
- Both child and parent have the same files open at the same position.
 - **But, since they are sharing file handles changes to the file offset made by the parent/child will be reflected in the child/parent!**

Calm Like A fork()bomb

What does this code do?

```
while (1) {
    fork();
}
```


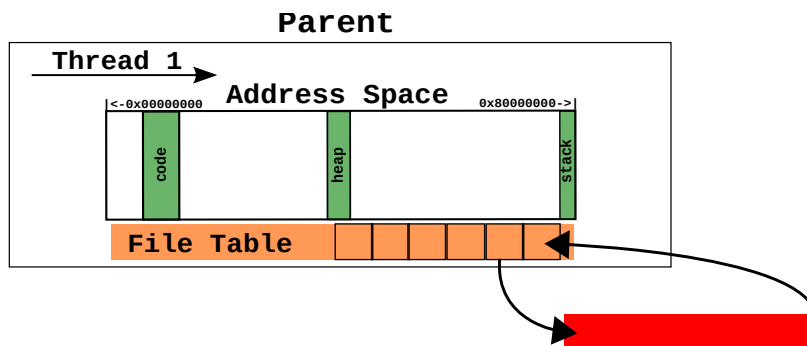

Pipes


Chains of communicating processes can be created by exploiting the pipe() system call.

- pipe() creates an anonymous pipe object and returns a two file descriptors: one for the read-only end, and the other for the write-only end.
- Anything written to the write-only end of the pipe is immediately available at the read-only end of the pipe.
- Pipe contents are buffered in memory.
- **Why is this useful?**

IPC Using fork() and pipe()

1. Before calling fork() the parent creates a pipe object by calling pipe().
2. Next, it calls fork().
3. After fork() the parent closes its copy of the read-only end and the child closes its copy of the write-only end.
4. Now the parent can pass information to the child.





```
# pipeEnds[0] gets the read end; pipeEnds[1] gets the write end.
int pipeEnds[2];

pipe(pipeEnds);

int returnCode = fork();

if (returnCode == 0) {

    # Don't need a loopback.
    close(pipeEnds[1]);

    # Read some data from the pipe.
    char data[14];
    read(pipeEnds[0], data, 14);
} else {

    # Don't need a loopback.
    close(pipeEnds[0]);

    # Write some data to the pipe.
    write(pipeEnds[1], "Hello, sweet child!\n", 14);
}
```

Issues with fork()

Copying all that state is expensive!

- Especially when the next thing that a process frequently does is start load a new binary which destroys most of the state fork() has carefully copied!

Several solutions to this problem:

- **Optimize existing semantics:** through copy-on-write, a clever memory-management optimization we will discuss in several weeks.
- **Change the semantics:** vfork(), which will fail if the child does anything other than immediately load a new executable.

- Does not copy the address space!

What if I don't want to copy all of my process state?

- fork() is now replaced by clone(), a more flexible primitive that enables more control:
 - over sharing, including sharing memory, and signal handlers,
 - and over child execution, which begins at a function pointer passed to the system call instead of resuming at the point where fork() was called.
- Try `man clone` in your CSE421 VM.

The Tree of Life

- fork() establishes a parent-child relationship between two process at the point when each one is created.
- The `ps` utility allows you to visualize these relationships.

```
challen@cse421:~$ pstree
init--apache2--apache2
                |--2*[apache2--17*[{apache2}]]
                |--2*[apache2--26*[{apache2}]]
                |--atd
                |--cron
                |--dbus-daemon
                |--dhclient3
                |--exim4
                |--6*[getty]
                |--git-daemon
                |--mailmanctl--8*[python]
                |--rsyslogd--3*[{rsyslogd}]
                |--sshd--sshd--sshd--bash--pstree
                |--udevd--2*[udevd]
                |--upstart-socket-
                |--upstart-udev-br
challen@cse421:~$
```